

PID discovery, verification, and durability blueprint

PrimeGaps lessons, semantic-transfer firewalls, a proposed exact-SxPID3 assurance architecture, autoresearch, and durable promotion

A source-pinned historical 20-lens review plus a 1 September 2026 70-lens publication-scope review (20 mandatory core and 50 additional artifact-specific lenses) and a separate 10-route comparison of a proposed exact categorical SxPID3 assurance architecture, formal-method limits, and durable repository practice.

2 September 2026

Contents

1	Status and scope	2
1.1	Publication contract	2
2	Executive decision	3
2.1	Adopt	3
2.2	Adopt conditionally	3
2.3	Reject or defer	3
3	Epistemic contract	4
3.1	What the tools can and cannot establish	4
4	Exact source inventory	4
5	What PrimeGapsLib proves	5
5.1	Mathematical skeleton	5
5.2	Discovery-to-theorem chain	6
5.3	Selected theorem and lemma map	6
6	The proof document and interactive graph	7
7	Repository and CI audit	7
8	Historical twenty-lens transfer council	7
8.1	Council synthesis and dissent	9
9	Dated 1 September 2026 adversarial publication closure	9
9.1	Twenty mandatory core lenses	9
9.2	Fifty additional artifact-specific hostile lenses	10
9.3	Ten materially distinct routes	12
10	The transfer ledger: adopt structure, not conclusions	13
10.1	Accepted architectural transfers	14
10.2	Conditional transfers	14
10.3	Rejected transfers	14
11	PID target: averaged empirical categorical SxPID3	15
12	Proposed proof-carrying semantic IR	17
12.1	View 1: semantic provenance graph	17
12.2	View 2: normalized exact arithmetic	17
12.3	Baseline and adoption gate	17
13	Layered assurance architecture	18
14	Formal methods as bounded instruments	18
14.1	Lean	19
14.2	SMT	19
14.3	Exact arithmetic and directed intervals	19
14.4	Exhaustive, mutation, metamorphic, and differential testing	19
14.5	Hashes, councils, and paper graphs	19
15	Attack matrix and required negative controls	19
16	Autoresearch without evidence laundering	20
16.1	Four distinct research loops	20
16.2	Promotion ledger	20
16.3	Council protocol	20
17	Repository durability and promotion	21
17.1	Why isolate work in branches and worktrees?	21
17.2	What failed, what detected it, and what repaired it	21
17.3	Formal verification and durable closure in one lane	22

17.4 Storage and trust ladder	24
17.5 Discover and freeze	24
17.6 Classify and preserve	25
17.7 Integrate	25
17.8 Verify and retire	25
17.9 Review-boundary safety disposition (dated 2026-09-01)	25
18 Implementation roadmap for SxPID3	25
18.1 Phase 0: preserve and adjudicate the packet	25
18.2 Phase 1: simple direct specification	26
18.3 Phase 2: dual formal semantics	26
18.4 Phase 3: exact products and intervals	26
18.5 Phase 4: keyed Rust refinement	26
18.6 Phase 5: bounded corpus and optional shards	26
18.7 Phase 6: promotion and external review	26
19 Evidence-card template	27
20 Final recommendation	27
21 Source-anchored claim register	27
22 Primary references	28

1 Status and scope

This document is a source-pinned research and engineering review, not a theorem and not an acceptance receipt. Its PrimeGapsLib observations are dated **19 August 2026**: it studies Axiom Math's PrimeGapsLib at exact Git commit 1faa7b14e82ddebc2772dfb9153922f01b106477, the interactive paper and its 132-page blueprint, then asks which proof-engineering patterns can improve pid-rs discovery and verification.

The current status of the proposed SxPID3 packet is governed separately by decision record 2 and the evidence-adjudication index, not by the dated PrimeGaps review. Decision record 2 keeps the complete target **proposed/open**; it credits two scoped results and leaves every load-bearing end-to-end cut open. The frozen revision-1 index remains historical intake evidence and is not silently rewritten as a current decision pointer.

The answer is deliberately asymmetric:

- several proof-engineering structures transfer usefully;
- the prime-gap mathematics does not transfer to partial information decomposition;
- an exact finite certificate does not validate a continuous estimator, a population model, a statistical procedure, or a scientific interpretation; and
- formal verification is one instrument in an argument, never a replacement for semantic review, mathematical judgment, falsification, numerical analysis, or independent custody.

The principal positive design is scoped to the proposed three-source empirical categorical shared-exclusions packet SX-CERTIFIED-AVERAGED-PID3-001. It introduces a typed semantic mass DAG, a separate reduced-rational product plan, forward and backward sparse-completeness checks, and a conditionally useful content-addressed sharded replay. This is a project-defined assurance architecture for an existing paper-defined functional. It is not a new PID measure or a currently accepted certificate.

The process is part of the result. This document therefore records the source inventory, formulas, theorem chain, the dated 20-lens transfer council, a separate 1 September 2026 70-lens adversarial closure (20 mandatory core plus 50 additional artifact-specific lenses), a separate 10-route comparison, dissent and limitations, the autoresearch boundary, the implementation roadmap, the evidence hierarchy, and the repository durability state machine. Exact path inventories, hashes, run rosters, and mutation records belong in machine-readable JSON or TSV companions rather than being copied repeatedly into human-facing migration PDFs.

1.1 Publication contract

This Markdown file is the canonical human narrative. The paired PDF is a derived publication, not a second editable source. Run `scripts/check-pid-discovery-verification-blueprint-pdf.sh` for the default committed-byte check, or `scripts/build-pid-discovery-verification-blueprint.sh` to regenerate it. The builder stages only this Markdown, its formatting filter/header, and four handcrafted SVG panels; canonicalizes the temporary root; fixes locale, time zone, creation epoch, document fonts, and temporary state; performs two isolated three-pass

LuaLaTeX builds; rejects compiler/layout warnings, non-A4 or non-PDF-1.7 output, missing semantic text anchors, and unembedded or non-Unicode-mapped fonts; and requires the two PDF byte strings to match. The title uses the repository-local lapis field, low-contrast girih-rosette pattern, saffron rule, turquoise subordinate accents, and an opacity fade. Decoration carries no claim or magnitude.

That equality is a **same-toolchain observation**, not a universal cross-platform theorem. A new Pandoc, LuaLaTeX, SVG renderer, font stack, or operating-system backend must rerun the gate and may legitimately produce different bytes. There is no accepted cross-toolchain equivalence relation or producer profile for this PDF: `--cross-toolchain` intentionally refuses rather than compare a convenient substitute. The four source SVG panels include titles, descriptions, numbered captions, and labels/patterns that do not rely on color alone. The PDF remains untagged; these source-level descriptions do not establish PDF /Alt, reading order, assistive-technology behavior, or PDF/UA conformance. The exact checker also renders every page at 120 dpi in color and grayscale, audits every PDF annotation and action, binds rendered HTTPS targets to the canonical source, and rejects relative/file links, undeclared web targets, actions outside the declared catalog/outline/link navigation whitelist, file specifications, and noncanonical annotation flags. This is bounded structural and rendering evidence, not accessibility or external link-availability evidence. The transfer ledger remains the machine-readable authority for the dated source hashes, anchors, semantic cards, and historical 20-lens registry. The adversarial publication review conducted on 1 September 2026 below is separate and does not rewrite that ledger. Neither review, the PDF, nor its polish upgrades the scientific status of the proposed SxPID3 packet.

Repository links remain relative in this canonical Markdown so forks and local checkouts stay portable. The PDF renderer rewrites only its declared repository-document links to canonical GitHub main-branch navigation URLs for this repository. Those mutable navigation URLs are not provenance; exact PrimeGaps source claims continue to use commit-pinned links.

2 Executive decision

2.1 Adopt

1. **Same-revision theory, certificate, and composite lanes.** Keep abstract mathematics, finite exact evidence, and the theorem that combines them separate, but require them to refer to the same immutable source revision.
2. **Untrusted discovery followed by a smaller checker.** Numerical search, autoresearch, and code generators may propose candidates; acceptance reconstructs the mathematical object from frozen input and checks it with a smaller, separately reviewed route.
3. **Explicit finite-to-semantic bridges.** A checked integer or rational inequality becomes useful only through named lemmas that connect stored data, mathematical definitions, and the final conditional claim.
4. **Assumptions in theorem and report types.** External premises, support conditions, units, estimand identity, and conditional status remain visible instead of being shortened away in summaries.
5. **Human-readable challenge plus machine comparator.** Freeze a small theorem roster, import and axiom policy, and an optional second-kernel replay independently of the candidate implementation.
6. **Typed dependency graph and completeness checks.** Sparse or packed evidence needs both directions: every required object is represented, and every represented object is semantically valid.
7. **Content-addressed evidence and promotion receipts.** If measured resource or recovery needs justify sharding a large finite replay, bind its partition and coverage as a certificate family rather than an opaque long job or expiring CI artifact. Shards are not part of the PID functional.

2.2 Adopt conditionally

1. Packing, custom kernels, and generated tables only after a simple streaming implementation is correct and profiling shows a real bottleneck.
2. A comparator only after its binary, version, theorem roster, allowed axioms, inputs, and full output are pinned and replayed in CI or under independently controlled custody.
3. A factor DAG inside SxPID3 revision 1 only if it leaves the frozen certificate semantics and resource admission policy unchanged. If the graph becomes an acceptance-bearing public object or changes an admission limit, the claim needs a new adjudicated revision.
4. Sharding only after the direct streaming baseline is correct and measured resource, replay, or recovery evidence justifies it. Internal evidence packaging may not change the accepted domain, result, or refusal policy; an acceptance-bearing shard schema or root requires an adjudicated claim/route revision.

2.3 Reject or defer

1. Direct transfer of GPY/Maynard sieve formulas, the number 246, the 50-tuple, factorial moments, or Bombieri-Vinogradov reasoning to PID.

2. Any implication from a finite exact categorical certificate to KSG consistency, continuous support, bias, variance, confidence coverage, population PID, causality, or scientific novelty.
3. Treating a green cached build, model consensus, self-assessed formalization, source hash, or polished blueprint as sufficient independent review.
4. Treating mutable blob/main source links or a status-colored dependency graph as a durable proof-to-paper correspondence receipt.
5. Kernel evaluation of millions of SxPID3 verdicts before a general proof and a bounded, independently replayable executable checker exist.

3 Epistemic contract

Every proposed transfer is represented by seven fields:

1. source meaning;
2. PID analogue;
3. assumptions preserved;
4. assumptions changed or introduced;
5. prohibited inference;
6. evidence required before adoption; and
7. disposition.

The complete machine-readable form is `audit/evidence/primegaps-to-pid-transfer-ledger-v1.json`. Shared words such as *information*, *lattice*, *factor*, *kernel*, *moment*, *certificate*, *positive*, or *independent* do not create a mapping theorem. Context determines meaning.

The council is advisory. Reviewers are asked to produce falsifiers and corrections, not votes. The integrator must reconstruct decisive claims from primary sources and owns the final scope. Two agents in one environment remain correlated evidence even when they use different roles or prompts.

3.1 What the tools can and cannot establish

Tool or method	Useful conclusion	Boundary that remains
Lean kernel	A term has the declared type under the recorded imports and axioms.	It does not prove that the formal statement faithfully represents the intended scientific claim.
Exact integer/rational checker	A finite algebraic predicate holds for reconstructed exact inputs.	It does not establish sampling validity, a population parameter, continuous support, or external provenance.
SMT solver	The bounded encoding is unsatisfiable or has a model under the encoded theory.	A missing premise, wrong encoding, or unsupported unbounded generalization remains possible.
Directed interval arithmetic	The declared exact real lies inside a rigorously constructed interval under its runtime assumptions.	It is not a confidence interval and does not establish model or estimator validity.
Mutation suite	Named faults are detected by the acceptance predicate.	It does not prove that all relevant faults have been enumerated.
CI	One bound revision passed one recorded environment and command family.	It does not prove universal portability, fresh reproducibility, or scientific correctness.
SHA-256	Two byte strings are equal except with the hash's cryptographic collision risk.	It does not establish truth, authorship, authenticity, semantic identity, or durable availability.
Multi-agent council	Diverse objections and candidate routes can be generated efficiently.	Shared training, prompts, sources, tools, custody, or objectives can create correlated blind spots.

4 Exact source inventory

The review uses the following identities. Reviewed means the bytes or Git object were inspected; it does not mean every statement was independently reproved.

Object	Exact identity	Review use
PrimeGapsLib tree	commit 1faa7b14e82ddeb2772dfb9153922f01b106477; tree 611f4b1caccb7850996cf425aefb1a6c42ed629f	Lean source, dependencies, comparator, history, and CI architecture
Interactive paper HTML	SHA-256 8ab349b59a897780ee7b023ec9b350442fc77b36aec84873019bfff2612a97c90; 10,793,548 bytes	statement/proof dependency graph, status, and source-link audit
Blueprint PDF	SHA-256 c93efd077ff4236926379d9879cc82e124a12f4eccc1c43a685054b6ba866117; 1,153,093 bytes; 132 A4 pages	human proof narrative and visual inspection
Blueprint layout text	SHA-256 1ea1f0fb79561a91edba050bc767325882aee83c51dd1ba169e2e0a5e25bd8cf; 495,392 bytes	searchable formula and statement review
Certificate source JSON	SHA-256 99808e1d95204c5bbc0cde845641cfdbb41ea950fa7fa9858df8bd093111787a; 55,303 bytes; 1,295 rows	untrusted candidate data input
Reviewed CI job log	SHA-256 408c6b2193069268c2250227e326d0aabd233244b1b6de804c9c98d4750c2b5; 33,294 bytes	observed warm-cache build behavior and skipped assurance routes

Primary locators are PrimeGapsLib, the interactive paper, and the blueprint PDF. Repository links in this review use the exact commit rather than mutable main.

The non-Git web/PDF/log rows above are **hash-only review observations**: their bytes are not shipped in this repository. Their digests neither provide a recovery route nor authenticate the publisher, and a future reacquisition may differ. The companion ledger marks that custody limit explicitly; the commit-pinned Git source is the only source family in this report with a direct object-recovery locator.

5 What PrimeGapsLib proves

The headline result is conditional. The paper states the result using the familiar liminf formulation

$$\liminf_{n \rightarrow \infty} (p_{n+1} - p_n) \leq 246.$$

whereas the reviewed public Lean declaration keeps the Bombieri--Vinogradov premise explicit and uses a frequently-at-infinity formulation:

```
bombieriVinogradov_implies_frequently_prime_gap_le_246
(hBV : BombieriVinogradov) :
exists frequently m atTop,
  nth Prime (m + 1) - nth Prime m <= 246
```

The two formulations belong to the same mathematical narrative, but this review does not silently substitute one theorem statement for the other. Neither may be summarized as an unconditional proof supplied by the certificate alone.

5.1 Mathematical skeleton

For an admissible tuple and nonnegative sieve weights, the analytic argument studies a difference of weighted sums of the form

$$S(N, \rho) = S_2(\lambda) - \rho S_1(\lambda).$$

The finite numerical route supplies a strict variational witness for

$$M_{k,\varepsilon} = \sup_F \frac{\sum_{m=1}^k J_{k,\varepsilon}^{(m)}(F)}{I_k(F)}.$$

At $k = 50$ and $\varepsilon = 1/25$, the repository represents a polynomial candidate as

$$F_c = \sum_i c_i b_{a_i, \alpha_i}$$

and checks, after clearing denominators, the exact inequality

$$4 \sum_{i,j} c_i c_j J_k^\varepsilon(i, j) < k \sum_{i,j} c_i c_j J_k^\varepsilon(i, j).$$

The exact Lean certificate conclusion used here is the strict threshold $M_{50,1/25} > 4$. The blueprint discusses a numerically constructed rational vector with an exact quotient above 4.0043, and a separate Lean file verifies that the parameter $Q = 4.004306$ makes a named positivity window nonempty. This review found no Lean theorem joining that stronger numeric margin to the stored certificate. It therefore does **not** describe the packed certificate as kernel-checking $M_{50,1/25} > 4.0043$.

Likewise, the blueprint's full degree-27 basis has 2,526 elements, while the reviewed JSON contains 1,295 nonzero sparse labels. The JSON row count is not the full basis size. The blueprint's degree-through-24 numerical observation is not presented here as a formal optimality theorem.

The numerical search that found the coefficient vector is not part of the trusted conclusion. The accepted route checks the exact integer/rational predicate, bridges the stored tables to the explicit quadratic forms, converts the explicit polynomial to an analytic L2 witness, derives $M_{\{50,1/25\}} > 4$, combines it with the analytic distribution premise, and uses an admissible 50-tuple whose diameter is checked to be 246.

5.2 Discovery-to-theorem chain

The source implements the following typed chain:

```

untrusted labelled JSON and search candidate
-> untrusted canonical generator
-> packed sparse tables
-> small Nat/Bool/RArray kernels
-> soundness bridges to direct mathematical recurrences
-> exact cleared-denominator inequality
-> explicit rational epsilon certificate
-> analytic L2 certificate
-> M_(50,1/25) > 4
-> Bombieri-Vinogradov and admissible H50
-> infinitely frequent prime gaps <= 246

```

The most transferable feature is not the particular recurrence. It is the refusal to trust the candidate generator, combined with local soundness lemmas that recover the mathematical specification from compact stored data.

5.3 Selected theorem and lemma map

Lean declaration	Role in the chain	What remains external or contextual
EpsCertificateExplicit	Finite rational coefficient family with a strict quadratic-form inequality.	The structure is specific to the sieve variational problem.
certFactorTables_correct	Connects cached factor tables to their declared mathematical factors.	It says nothing about PID event masses or lattices.
certStoredInequality	Kernel-decided strict inequality over stored row totals.	Stored-row correspondence needs separate lemmas.
certLhsStoredRowSum_eq, certRhsStoredRowSum_eq, certInequality	Recover the original left and right quadratic forms from stored rows. Produces the complete cleared-denominator witnessing inequality.	They share the certificate interpretation and source data. It is still a finite certificate predicate.
EpsCertificateExplicitPrecomputedInt.toExplicit	Bridges packed exact data and direct moments to the explicit rational certificate.	Integer runtime and soundness premises remain in the trust analysis.
EpsCertificateExplicit.toAbstract	Converts the explicit rational polynomial into the analytic L2 certificate.	This bridge is prime-gap-specific; only its architectural role transfers.
EpsCertificate.four_lt_MEps	Derives the strict variational inequality from the analytic witness.	It does not provide the Bombieri-Vinogradov premise.
existsEpsCert50	Packages the checked certificate as the endgame hypothesis.	The existential is only one premise of the final theorem.
lem_theta_bv	Uses Bombieri-Vinogradov to choose a certificate-admissible distribution level.	This is the analytic black box of the 246 development.
card_H50, diameter_H50, admissible_H50	Kernel-check the explicit 50-tuple, its diameter, and admissibility.	They do not prove the comment's "narrowest" optimality claim; the tuple has no PID analogue.
thm_main	Combines the analytic and finite routes to obtain frequent gaps at most 246.	The conclusion remains conditional on BombieriVinogradov.

Exact source anchors are retained in the transfer ledger. The theorem table is an architectural map, not an independent proof audit of every imported lemma.

6 The proof document and interactive graph

The blueprint is a readable 132-page mathematical document. Visual inspection of every page at contact-sheet scale and representative pages at higher resolution found consistent margins, legible formulas, stable headers, and useful per-statement Lean annotations. This is a substantial review aid.

The interactive paper contains 255 typed entities: 126 marked formalized, 21 marked gap, and 108 marked unknown. The graph records 1,126 typed uses: 805 statement uses and 321 proof uses. The observed dependency edges are acyclic and backward-pointing, with a maximum depth of 27 to the final theorem. This makes the site useful for locating high-centrality definitions and incomplete cuts.

The graph is not, in its present form, a durable correspondence certificate:

- 186 observed repository source links use mutable blob/main URLs;
- status and link presence are not equivalent: some formalized or gap nodes lack links, while some unknown nodes have links;
- later paper sections are largely marked unknown despite substantial Lean source elsewhere in the repository;
- PaperTag.lean constructs its blueprint URL from a literal todo string; and
- its duplicate-handling path logs “skipping duplicate” but still inserts the duplicate entry; and
- no same-revision receipt binds the HTML, PDF, Lean tree, graph edges, statuses, and source spans.

The concept transfers: maintain a revision-pinned claim/proof/evidence DAG. The current mutable site implementation does not transfer as an assurance mechanism.

7 Repository and CI audit

PrimeGapsLib separates PrimeGapsTheory, PrimeGapsCert, and their composite library. That split is valuable because a reviewer can build or inspect the mathematical theory without always replaying the large computation.

The reviewed repository nonetheless has important reproducibility and custody limits:

- the reviewed commit's reachable history contains seven commits and begins with a large monolithic import; no release tag or GitHub release asset is bound into this report;
- PrimeNumberTheoremAnd is requested from mutable main, although the manifest records the resolved commit used by this checkout;
- the certificate JSON is transitively bound by the reviewed Git tree, but has no explicit sidecar content digest or candidate-generation/search receipt beyond that repository-object binding;
- the workflow uses mutable action version tags rather than full commit pins;
- the latest reviewed run restored roughly 2.47 GB of cache, built successfully, but skipped the comparator, leanchecker, and nanoda routes and uploaded no proof receipt;
- the comparator configurations are documented but not wired into the reviewed CI workflow; and
- the reviewed project metadata describes self-assessed and internal review, and this report binds no independent human or institutional review receipt.

None of these observations refutes the Lean theorem. They limit claims about fresh reproducibility, external verification, review independence, and durable publication.

One parser detail illustrates why even a good architecture cannot be copied mechanically. The PrimeGaps source generator converts some labelled JSON integers using `Int.toNat`, so a negative input can be coerced rather than rejected. For an existential witness theorem this can be safe: the downstream checker only needs the reconstructed object to be *some* valid witness. A PID report, by contrast, is a deterministic function of the exact experimental count table. Its decoder must reject such coercions, require canonical bytes, and bind the accepted data identity.

8 Historical twenty-lens transfer council

This dated council belongs to the 19 August PrimeGaps review. It is a structured attack, not a vote. Each lens asks a different question of the same proposed transfer. Agreement among lenses does not create evidence, and a single load-bearing objection can veto promotion. The detailed evidence and nonimplications are preserved in the machine-readable transfer ledger. The table below uses the same canonical names, numbering, and order as that historical 20-lens registry; it is the human synthesis of those exact rows. The separate current closure in the next section does not retroactively change these rows.

#	Lens	PrimeGaps observation	PID decision
1	Estimand	The public result concerns bounded gaps between primes.	Name the deterministic averaged empirical categorical SxPID3 functional; transfer proof roles only, never sieve objects or constants.
2	Types	One fixed existential rational polynomial witness suffices for the source theorem.	PID accepts bounded runtime count tables, so canonical bytes, row roles, semantic masses, exact plans, intervals, and dispositions require distinct types.
3	Quantifiers	The public theorem retains an explicit Bombieri-Vinogradov premise and an existential certificate.	State the finite-input SxPID3 implication and its quantifier order; do not replace it with population, estimator, asymptotic, or causal prose.
4	Mappings	No theorem maps GPY/Maynard objects to shared-exclusions PID.	Transfer assurance roles only; any mathematical value or proposition needs a positive premise-checked target map.
5	Support/reference measure	PrimeGaps' analytic measure premises are internal to its simplex and sieve problem.	Reconstruct positive empirical support from the PID definition; do not infer continuous support, a reference measure, or KSG validity.
6	Lattice/source count	Packed signatures and recurrence indices form source-specific carriers.	Freeze the 18-node three-source antichain poset, its order, event semantics, and Möbius direction independently.
7	Units/gauge	The certificate compares dimensionless exact rational forms.	Preserve natural logarithms, nats, informative/misinformative/net orientation, normalization, and any later preprocessing gauge.
8	Population/empirical	The source certificate is an exact finite witness inside a number-theory proof, not a statistical estimate.	Keep empirical PMF functional, estimator output, fixture, represented binary value, and population quantity separate.
9	Sampling	The fixed prime-gap witness does not supply PID row-law assumptions.	The narrow exact categorical claim is deterministic; any later population, bootstrap, permutation, or confidence statement needs its own sampling contract.
10	Selection/UQ	Numerical search found a successful candidate, but the reviewed project lacks a complete exposure, stopping, and rejected-candidate ledger.	Freeze objectives and judges before candidate access, retain adaptation history, and never treat rediscovery or consensus as calibration.
11	Formal correspondence	Lean bridges explicit objects, but the website graph has mutable links and status inconsistencies.	Bind a reviewed paper-to-formula map for event, averaging, order, units, and conclusion before formal code earns semantic credit.
12	Kernel/axioms/toolchain	The project pins Lean and resolved dependencies, while reviewed CI skips comparator, leanchecker, and nanoda.	Inventory kernel/runtime/import closure/axioms and require fresh fail-closed replay; sorry, theorem premises, conventional axioms, and noncomputable are different categories.
13	Route dependency diversity	Theory, certificate, packed kernels, and optional comparator diversify some cuts but share source mathematics and data.	Record shared semantic, parser, arithmetic, custody, and model cuts; route or agent count is not independence.
14	Numerical/binary64	Approximate search is replaced by an exact rational acceptance inequality.	Use exact products for finite sign and separately verified directed intervals for log magnitude; neither validates binary64 estimators or platform libm.
15	Compiled/wrapper parity	The reviewed certificate route does not exercise pid-rs Rust/Python surfaces.	Compare all 108 keyed scalar audit expressions across specialized/general, serial/parallel, feature, build, platform, serialization, and Python routes.
16	Counterexample/mutation	Small exact kernels are valuable only if load-bearing premises actually affect acceptance.	Mutate decoding, event connectives, target restriction, count weighting, lattice order, products, intervals, resources, claims, and coordinated reseals.
17	Custody/threat/refusal	Public source and a warm-cache build do not provide an external-custody comparator or certificate receipt.	Freeze candidate scope and judge authority, refuse unavailable verification, and retain exact acquisition/execution records plus same-actor limits.
18	Citations/novelty	PrimeGaps cites its mathematical sources and presents a formalization.	Pin primary statements and application hypotheses; classify the PID assurance design as project-defined engineering, not a new measure or theorem.
19	Ecosystem/authority	The interactive graph is useful navigation metadata, not an acceptance authority or proof-term DAG.	Bind claim, source, theorem, executable, consumer, and archive authority directions at one revision.
20	Resource/platform/release	Packed computation is fast in a warm cached run, but no reviewed same-revision release receipt closes the lane.	Start with a bounded streaming route; freeze resource refusal, attempts, withdrawals, OS/toolchain scope, release assets, durability, and recovery before promotion.

8.1 Council synthesis and dissent

- **Strongest surviving mechanism:** untrusted candidate generation followed by exact reconstruction in a smaller checker and a named semantic bridge.
- **Strongest objection:** identical words such as *factor*, *certificate*, *positive*, or *lattice* can conceal different objects and different sufficient conditions.
- **Most likely false assumption:** that two nominally independent routes do not share the same paper transcription, generated registry, parser, arithmetic library, model lineage, or custody.
- **Evidence that would change the decision:** a pinned primary-source map; a generated-answer-blind direct exact implementation; concrete Lean and separately encoded finite-SMT routes; complete hostile models; dependency-disjoint interval reconstruction; keyed Rust refinement; and a fresh, externally held replay.
- **Dissent retained:** the semantic DAG may be unnecessary. A simpler direct reconstruction is the baseline and remains preferred unless profiling and proof size show a measured benefit.

9 Dated 1 September 2026 adversarial publication closure

This is a new review of the present publication and decision boundary, not an amendment to the historical 20-lens ledger. Twenty mandatory core lenses and fifty additional artifact-specific lenses were applied beyond the ordinary author pass: seventy typed rows in total, with each row counted once. Thematic overlap between a core question and a finer artifact-specific attack is intentional; the count is coverage, not independence. Ten materially distinct routes were then compared separately, including no action and the simplest applicable non-PID route. The counts are coverage checklists, not mathematical evidence. The closure verdict is narrow: publish this document as a proposed/open research blueprint, preserve the two scoped decision-record-2 results, and grant no complete-certificate, estimator, population, application, novelty, or cross-toolchain-PDF credit.

The audit subject is this blueprint's current wording and the open certificate boundary that it reports. Each row below has exactly one typed disposition. PASS identifies current-byte evidence; CORRECT records an old and corrected statement; NARROW retains an explicit scope; OPEN names an unclosed obligation; NEGATIVE retains a counterexample or prohibited route; and N/A gives a typed applicability reason. A publication-level pass never closes the underlying certificate.

9.1 Twenty mandatory core lenses

Each mandatory core row has one disposition and scoped evidence. Where a core lens contains several subquestions, the disposition reports the least favorable load-bearing state instead of averaging or selecting a favorable subpart.

#	Mandatory core lens	Typed disposition and scoped evidence
1	Estimand	NARROW - The target is only the averaged empirical categorical Makkeh--Gutknecht--Wibral SxPID3 functional, with event formation, count weighting, logarithm, and Möbius inversion ordered explicitly; no estimator, population, or other-PID estimand is included.
2	Types	NARROW - The source/target alphabets, nonnegative integer count table, supported rows, antichain carrier, exact rational products, real-valued nats, and 108 keyed audit outputs are typed; pointwise, continuous, fitted, and application outputs remain excluded.
3	Quantifiers	NARROW - The formulas quantify over every supported row and all 18 frozen antichains for one finite table; the bounded corpus quantifies only over binary count vectors with $1 \leq N \leq 5$, with no limit, probability-over-samples, or arbitrary-alphabet claim.
4	Mappings	OPEN - No shared vocabulary is treated as a mapping theorem. The primary-paper-to-formula map and complete semantic bridge remain Program A obligations; the PrimeGaps mathematics and other PID definitions are explicitly not transferred.
5	Support/reference measure	NARROW - The target uses a finite empirical PMF on positive-count support relative to the finite counting reference; density, Radon--Nikodým, singular-support, and continuous-boundary premises are not part of this estimand and receive no transfer credit.
6	Lattice/source count	OPEN - Three ordered sources, 18 nonempty antichains, the order, zeta orientation, and Möbius direction are frozen, but carrier generation, complete order derivation, event semantics, and paper correspondence remain open.
7	Units/gauge	NARROW - Natural logarithms determine nats and $1/N$ performs empirical averaging. A paper reference expressed in bits is converted explicitly by multiplication with $\ln 2$, never silently relabeled; no metric, neighbourhood radius, preprocessing scale, or continuous gauge is imported.
8	Population/empirical	NARROW - The document separates one empirical PMF, its exact functional value, bounded fixtures, and prospective executable reports from any population law or estimator output.
9	Sampling	N/A - Row-law, i.i.d., stationary, ergodic, dependence-rate, split, and block-length assumptions are not needed to evaluate the deterministic finite-table functional; any inferential wrapper must create a separately adjudicated claim.

#	Mandatory core lens	Typed disposition and scoped evidence
10	Selection/UQ	OPEN - Statistical null, confidence, and multiplicity semantics are not applicable to the deterministic exact functional, but adaptive candidate selection is live: its exposure, withdrawal, stopping, held-out judge, and first-result receipt remain required before promotion.
11	Formal correspondence	OPEN - The concrete Lean/SMT propositions, source-to-symbol map, and finite-to-semantic bridges remain incomplete; generic algebra or a few fixtures do not establish the intended paper functional.
12	Kernel/axioms/toolchain	OPEN - Exact proposition/import/axiom/toolchain rosters and fail-closed concrete Lean plus solver-neutral SMT replays are prospective; no unavailable verifier may be treated as success.
13	Route dependency diversity	OPEN - The route comparison records generic shared paper, registry, parser, host, toolchain, data, reviewer/model-lineage, and custody cuts, but no complete per-route-pair functional/epistemic/institutional independence vector or institutional replay exists.
14	Numerical/binary64	OPEN - Exact rational sign/zero has bounded partial evidence, while directed magnitude intervals, aggregate error budgets, overflow/refusal behavior, binary64/libm/compiler effects, and compiled refinement remain open.
15	Compiled/wrapper parity	OPEN - Stable-key Rust traces, specialized/general parity, serial/parallel and build-feature matrices, Python bindings, serialization, and failure surfaces are not yet acceptance evidence.
16	Counterexample/mutation	OPEN - Exact zero and negative-atom witnesses plus named semantic attacks are retained, but the complete causal mutation suite and fault-model coverage remain unclosed.
17	Custody/threat/refusal	OPEN - Local source-bound receipts and refusal rules exist, but candidate/judge authority, complete threat records, independently acquired replay, external custody, and terminal atomic-install evidence remain incomplete.
18	Citations/novelty	OPEN - PrimeGaps claims are commit-pinned and the PID/Möbius references plus project-defined no-priority boundary are identified, but the exact version-pinned MGW statement-to-formula map and imported hypotheses remain Program A obligations.
19	Ecosystem/authority	OPEN - Consumer deployment and application authority are not claimed, but capability gaps, proposed/open status, wrapper authority, and current Markdown-to-PDF projection freshness remain live obligations; a non-PID alternative check precedes any later application claim.
20	Resource/platform/release	OPEN - Preflight, cancellation, determinism, platform/toolchain observations, exact release identity, hosted execution, recovery, attempt/withdrawal disposition, and retirement evidence are not all closed.

9.2 Fifty additional artifact-specific hostile lenses

#	Hostile lens	Typed disposition and current-byte evidence or retained boundary
1	Scientific-object identity	NARROW - The status and target sections restrict the object to the averaged empirical categorical SxPID3 functional on one finite count table, not a population parameter or estimator output.
2	PID-functional identity	PASS - The target section names Makkeh--Gutknecht--Wibral shared exclusions and explicitly excludes Williams--Beer, BROJA, continuous Ehrlich shared exclusions, and every target-free analogue.
3	Variable-role identity	PASS - The displayed row and event formulas give three ordered sources and one target different roles; no relabeling may silently exchange a source with the target.
4	Alphabet and support	NARROW - The equation preamble requires declared finite alphabets and nonempty finite empirical support; the bounded audit further fixes all four alphabets as binary.
5	Empirical versus population	NARROW - The equation preamble and nonclaims retain deterministic empirical-PMF scope; exact count-table evaluation supplies no sampling, consistency, bias, variance, confidence, or causal statement.
6	Sampling and dependence	N/A - Independent-row, exchangeability, stationarity, and block-length premises are not required to evaluate this deterministic functional; any inferential wrapper must add and justify its own premises.
7	Source-event connective	PASS - The displayed event is a union across antichain branches and an intersection within each branch; the attack matrix requires an exact witness for either reversal.
8	Target restriction	PASS - The displayed misinformative mass intersects the source event with the realized target event; the attack matrix requires an exact witness for omission or relocation of that restriction.
9	Antichain semantics	PASS - The target section defines every node as a nonempty family of nonempty source subsets with no two distinct members comparable by inclusion.
10	Carrier completeness	OPEN - The frozen list has 18 keys, but generation, uniqueness, exclusion of the zero mask, and paper correspondence remain unclosed obligations; a hard-coded list is not a proof.
11	Source-count boundary	NARROW - Eighteen is retained only as the three-source carrier size; 166 is the separate four-source carrier and is outside this audit.
12	Redundancy-order direction	OPEN - The quantified subset direction is frozen and displayed, but its complete source-derived semantic justification remains in Program A; a reversed order can be coherently wrong.

#	Hostile lens	Typed disposition and current-byte evidence or retained boundary
13	Zeta orientation	NARROW - The frozen convention and displayed equations fix cumulatives as rows and atoms as columns; the exact matrix facts do not close paper correspondence.
14	Möbius direction	NARROW - The displayed inverse maps cumulative vectors to atom vectors, while the text explicitly withholds semantic credit from algebraic invertibility alone.
15	Component orientation	NARROW - The formulas keep informative, misinformative, and signed-net components separate, with signed net equal to informative minus misinformative; the paper-to-formula bridge remains open.
16	Count weighting	PASS - The cumulative products use exponent c_z and the values use factor $1/N$, rather than one vote per distinct supported row.
17	Units and logarithm base	PASS - The natural logarithm fixes the unit as nats, while $1/N$ performs empirical averaging; the formulas keep these distinct roles, and prohibit bit labels or omitted averaging.
18	Exact-zero meaning	PASS - The normalized-rational section gives $(2 \cdot 6)/(3 \cdot 4) = 1$ as a retained counterexample to structural or empty-plan zero tests.
19	Negative-value preservation	PASS - The target section gives the exact bounded witness $N = 3$, key 02+04, and $R^{sx} = 9/16 < 1$; it and the attack matrix prohibit clamping, absolute values, or positive-only summaries.
20	Coordinate ontology	CORRECT - Old shorthand could be read as “108 PID coordinates”; the corrected statement is 108 keyed scalar audit expressions over 18 lattice positions, not 108 atoms, nodes, or independent parameters.
21	Algebraic dependence	PASS - The target section records net blocks as derived from component blocks and atom blocks as fixed transforms of cumulative blocks; coverage count is not dimension count.
22	Source permutation	OPEN - The semantic-IR obligations and attack matrix require simultaneous transport of alphabets, cells, events, keys, and output identities; no completed keyed compiled proof is claimed.
23	Integer replication	NARROW - Multiplying all counts by one positive integer preserves the represented empirical law; it is an executable regression only when the rescaled input remains admitted by the declared domain/resource contract, not a new theorem or distinct-law count.
24	Domain extrapolation	NARROW - The bounded result is explicitly limited to binary alphabets and totals one through five; it grants no arbitrary-total, arbitrary-alphabet, full-support, or pointwise credit.
25	Canonical parsing	OPEN - D1 remains open: accepted bytes still need one strict decoding, member-order policy, integer grammar, unique reserialization, and rejection of noncanonical alternatives.
26	Integer coercion	NEGATIVE - The retained cross-context comparison shows that negative-to-natural coercion can be non-load-bearing only in a separately validated existential candidate lane; it is prohibited for a deterministic PID count table.
27	Duplicate and unknown fields	OPEN - D1 must reject duplicate members, missing keys, unknown acceptance-bearing fields, reordered identities, and noncanonical encodings rather than normalize them silently.
28	Data and software identity	PASS - The epistemic and evidence-card sections keep input bytes, source revision, executable identity, build context, result identity, truth, and authenticity as separate fields.
29	Resource preflight	OPEN - The architecture requires aggregate and per-coordinate checks before expensive powers or allocations; the complete SxPID3 admission implementation is not yet accepted.
30	Cancellation and atomicity	OPEN - The roadmap makes cancellation, timeout, precision exhaustion, and storage failure terminal, but the complete install/refusal path remains prospective.
31	Rational normalization	NARROW - Exact sign authority is the normalized positive rational compared with one; this algebraic fact does not prove correct decoding, events, or executable refinement.
32	Semantic-factor aliasing	NEGATIVE - Equal-valued target and event masses remain a retained counterexample to value-only node merging until a proved evaluation map reaches the arithmetic plan.
33	Directed-log enclosure	OPEN - Nonzero magnitudes still need independently reconstructed outward enclosures under recorded precision, backend, range-reduction, and tail-bound contracts.
34	Interval decision rule	OPEN - The proposed verifier requires containment rather than overlap; an interval touching zero has no sign/zero authority because the exact-product lane decides, while producer/verifier implementation and dependency separation remain unclosed.
35	Binary64 boundary	NARROW - Exact rational sign and real-log enclosure are expressly withheld from floating-point KSG, libm, or current Rust validation.
36	Compiled-route refinement	OPEN - Specialized and general Rust paths still need stable-key value traces from bound binaries; the credited lexical routing observation computed no values.
37	Build and parallel parity	OPEN - Serial/parallel, debug/release, feature, platform, and toolchain comparisons remain separate required finite observations, not portable identity theorems.
38	Binding and serialization	OPEN - Python exposure, report encoding, field order, and stable keys require their own refinement checks; a binding is not method-origin evidence.
39	Source-to-formula correspondence	OPEN - Program A still needs a pinned primary-paper map for event, averaging, order, units, and conclusion before formal or executable evidence earns semantic credit.
40	Lean trust boundary	NARROW - Kernel credit is restricted to the exact proposition, imports, axioms, generated code, and recorded holes; it grants no source-correspondence or runtime claim.
41	SMT trust boundary	NARROW - Unsatisfiability concerns only the encoded finite theory; unknown is no evidence and two solvers can share the same mistranscription.
42	Mutation causality	OPEN - Each accepted hostile mutation must fail through its intended predicate, but the complete causal suite and fault-model coverage are not claimed.

#	Hostile lens	Typed disposition and current-byte evidence or retained boundary
43	Route independence	NARROW - The route section and shared-cut paragraph distinguish language, algorithm, parser, data, semantic, toolchain, host, model-lineage, reviewer, and custody cuts; shared source and local custody remain explicit.
44	Current-decision freshness	CORRECT - Old wording treated the 31 August closure and Program rows as current; the corrected 1 September closure binds the current decision/index bytes, which retain D1 and every Program A--E status, while coordinated resealing is not an authenticity theorem.
45	Selection and autoresearch	OPEN - The promotion ledger freezes grammar, objective, budget, exposure, stopping, withdrawals, and judge authority, but it supplies no acceptance receipt for a future adaptive run.
46	External custody	OPEN - Same-repository hashes and same-host reruns provide no independent acquisition, authenticity, or institutional-review receipt; fresh separately held replay remains required.
47	Link and PDF portability	CORRECT - Old prose said all active actions were rejected; the corrected contract admits only the declared catalog/outline/link navigation whitelist, keeps Markdown links relative, and limits exact rendering to the same toolchain.
48	Durability and recovery	CORRECT - Old “remote proof” shorthand is replaced by remote-ref, ancestry, hosted-run, and recovery-drill checks; local refs, remote main, releases, and archives remain different states, and retirement needs a tested successor.
49	Application and non-PID alternative	NARROW - PID is retained only when it answers a justified decomposition question; MI, CMI, direct task loss, ablation, coverage, or explicit failure utility veto forced PID use.
50	Provenance, novelty, and citation	CORRECT - The private design-source name was removed in favor of the repository-local publication adaptation; the target definition and classical Möbius inversion are cited, and the assurance composition is project-defined with no priority claim.

9.3 Ten materially distinct routes

#	Route	Decision, reason, and reconsideration evidence
1	No action; preserve the dated packet only	Retained as the safe fallback if source correspondence stays ambiguous. Rejected as the active publication route because it leaves current status and useful bounded evidence difficult to find.
2	Choose the simplest non-PID application analysis and withdraw only the application-specific PID claim	Selected whenever the scientific question is answered by MI, CMI, direct task loss, ablation, coverage, or failure utility. It does not withdraw or certify the separate repository functional-certification packet; it prevents method forcing in that application.
3	Prose review plus a few hand fixtures	Useful for discovery and teaching, but rejected as acceptance evidence because it cannot establish the complete carrier, registry, parser, or bounded corpus.
4	Treat the two existing Python routes and matching digest as final	Rejected as complete certification because both routes share the publication transcription, registry, runtime class, host, framing, and local custody. Their bounded result remains credited.
5	Simple generated-answer-blind direct exact checker	Selected as the baseline. It must reconstruct canonical counts, events, order, matrices, products, and keyed signs without importing generated answer tables. It still depends on a separately reviewed primary-source specification.
6	Semantic mass DAG plus separate reduced-rational plan	Selected conditionally as an internal refinement only if a proof of evaluation, forward and backward completeness, mutations, and benchmarks beat the direct baseline without changing the claim.
7	Lean-only concrete formalization	Selected as one supporting route, never as the whole argument. It can close concrete finite semantics but not source correspondence, runtime bytes, intervals, Rust refinement, or custody.
8	SMT-only finite encoding	Selected as a differently encoded supporting route under at least two solvers. It remains subordinate to a raw constraint derivation, satisfiable hostile models, and explicit unknown handling.
9	Immediate custom packing or mandatory content-addressed sharding	Deferred. The unsharded streaming route remains authoritative; packing or internal evidence shards earn adoption only after it is correct and measured resource, replay, or recovery evidence shows a real need. An acceptance-bearing shard contract requires a new adjudicated revision.
10	Layered composite with direct baseline, conditional IR, dual formal routes, intervals, compiled refinement, optional shards, and external replay	Selected overall because each layer owns a different cut and no downstream layer repairs an upstream semantic error. Promotion still requires every frozen obligation, including upstream canonical-decoder obligation D1, every Program A--E closure condition, and fresh external custody.

The selected design takes the baseline from route 5, the provenance separation from route 6, the bounded formal diversity from routes 7 and 8, the measured-only optimization rule from route 9, and the end-to-end discipline from route 10. Route 2 remains an application-level veto: even a perfect PID certifier is irrelevant when PID adds no justified information to the real decision.

Shared failure cuts remain the paper transcription, the exact registry, parser semantics, one host and toolchain family, common data, same-actor custody, and correlated model or reviewer blind spots. Dissent is preserved: the semantic DAG may add more proof surface than value; exact sharding may be premature; and the correct application decision can be to use no PID at all. New evidence can change those dispositions only through a premise-explicit source map, generated-answer-blind direct replay, measured resource results, causal hostile tests, and fresh independently held execution.


10 The transfer ledger: adopt structure, not conclusions

Every row in the companion ledger has source semantics, a PID analogue, preserved assumptions, changed assumptions, prohibited inferences, required evidence, and a disposition. The decisions compress to the following map.

FIGURE 1 - TRANSFER FIREWALL

Begin with source meaning, not shared vocabulary

The first two panels record the source object and the required mapping discipline.



Panel labels and hatch/dot patterns supplement the publication palette; they do not assert PDF accessibility conformance.

Figure 1: Semantic transfer firewall, part 1: source meaning and the required transfer card.

FIGURE 2 - TRANSFER FIREWALL

Adopt a proof role only after PID-specific evidence

The proposed assurance route is not a new PID measure and does not transfer the source conclusion.




Figure 2: Semantic transfer firewall, part 2: the PID-specific assurance route and nontransfer boundary.

10.1 Accepted architectural transfers

1. **Theory / certificate / composite separation.** The theory lane states general categorical semantics. The certificate lane checks one finite input. The composite lane is the only lane allowed to state the bound end-to-end implication.
2. **Untrusted discovery / small checker.** Search and generation remain replaceable. Acceptance reconstructs the object from canonical bytes rather than trusting a producer's tables.
3. **Finite-to-semantic bridges.** Stored products, matrices, or graph nodes earn meaning only through checked evaluation and refinement statements.
4. **Premises in types and reports.** Theorem inputs and report fields preserve units, support, estimand, source, and conditional status.
5. **Content-addressed promotion.** Exact inputs, outputs, tools, and evidence move together under one same-revision manifest.

10.2 Conditional transfers

- A human challenge and comparator are useful only when executed with pinned binaries, theorem and import rosters, axiom policy, raw results, and hostile tests.
- A proof-document DAG is useful only when revision-pinned and checked for unique identities, valid spans, status coherence, and evidence direction.
- Sparse graphs, packing, and sharding are implementation choices. They require completeness proofs and benchmarks; they are not mathematical evidence by themselves. An acceptance-bearing representation or admission-policy change requires a new adjudicated claim/route revision.
- Autoresearch may adaptively search, but a candidate-inaccessible judge and a complete exposure ledger control promotion.

10.3 Rejected transfers

- No GPY or Maynard expression, eigenvector objective, admissible tuple, or number 246 transfers to PID without a new theorem establishing an actual typed map.

- No exact categorical product validates KSG, continuous shared exclusions, population PID, or a statistical procedure.
- No warm-cache build, self-review, agent consensus, mutable website, or polished PDF supplies independent acceptance.

11 PID target: averaged empirical categorical SxPID3

The chosen pilot is narrow because it has exact finite semantics and a real open assurance packet. It is the categorical Makkeh--Gutknecht--Wibral shared-exclusions functional. It is **not** Williams--Beer $I_{\min}$, BROJA, the continuous Ehrlich shared-exclusions construction, a target-free analogue, the full mixed-dimensional continuous PID3, KSG, fitted quantization, or pointwise certification. These measures and estimands have different definitions and assumptions; their atoms must not be pooled or relabeled. The complete packet remains proposed/open. Current decision-v2 evidence is partial and bounded; it does not close any end-to-end implication.

The assumptions for every equation in this section are: three ordered finite source alphabets $\mathcal{S}_1, \mathcal{S}_2, \mathcal{S}_3$; one finite target alphabet $\mathcal{T}$; one finite, nonempty empirical count table with nonnegative integer cell counts; natural logarithms and nats; and the deterministic empirical-PMF estimand. Zero-count cells are omitted from the positive support. No equation in this section asserts a population, sampling, continuous-support, estimator, calibration, or application result. The bounded audit later adds the stricter assumptions that every alphabet is binary and $1 \leq N \leq 5$.

Let

$$\mathcal{Z} = \mathcal{S}_1 \times \mathcal{S}_2 \times \mathcal{S}_3 \times \mathcal{T}, \quad c_z \in \mathbb{N}_0, \quad \mathcal{Z}_+ = \{z \in \mathcal{Z} : c_z > 0\}, \quad N = \sum_{z \in \mathcal{Z}} c_z > 0.$$

A supported joint realization is

$$z = (s_1, s_2, s_3, t) \in \mathcal{Z}_+.$$

An antichain α is a nonempty family of nonempty subsets of $\{1, 2, 3\}$ in which no two distinct members are comparable by inclusion. For the keyed row z , define the source event

$$E_\alpha(z) = \bigcup_{a \in \alpha} \bigcap_{i \in a} \{S_i = s_i\}.$$

The union across antichain branches and conjunction within a branch are load-bearing. Write the target event as $T(z) = \{T = t\}$. Define

$$\begin{aligned} U_{\alpha,z} &= \sum_{z' \in E_\alpha(z)} c_{z'}, \\ V_{\alpha,z} &= \sum_{z' \in E_\alpha(z) \cap T(z)} c_{z'}, \\ T_z &= \sum_{z' \in T(z)} c_{z'}. \end{aligned}$$

For valid supported rows these masses are positive. The informative, misinformative, and net local ratios are

$$\rho_{\alpha,z}^+ = \frac{N}{U_{\alpha,z}}, \quad \rho_{\alpha,z}^- = \frac{T_z}{V_{\alpha,z}}, \quad \rho_{\alpha,z}^{\text{sx}} = \frac{\rho_{\alpha,z}^+}{\rho_{\alpha,z}^-} = \frac{NV_{\alpha,z}}{U_{\alpha,z}T_z}.$$

Empirical weighting is by c_z/N , not uniformly over distinct supported rows. Clearing the logarithmic sums into exact positive-rational products gives

$$\begin{aligned}
Q_{\alpha}^{+} &= \prod_{z \in \mathcal{Z}_{+}} \left(\frac{N}{U_{\alpha,z}} \right)^{c_z}, \\
Q_{\alpha}^{-} &= \prod_{z \in \mathcal{Z}_{+}} \left(\frac{T_z}{V_{\alpha,z}} \right)^{c_z}, \\
Q_{\alpha}^{\text{sx}} &= Q_{\alpha}^{+}/Q_{\alpha}^{-}, \\
C_{\alpha}^u &= \frac{1}{N} \ln Q_{\alpha}^u, \quad u \in \{+, -, \text{sx}\}.
\end{aligned}$$

There are exactly 18 nonempty antichains of the seven nonempty subsets of three sources in the frozen registry. The complete carrier, stable keys, zeta rows, sparse inverse, and six-block coordinate order are listed in the frozen conventions. The redundancy order is

$$\alpha \preceq \beta \iff \forall b \in \beta \exists a \in \alpha : a \subseteq b.$$

With cumulative rows and atom columns,

$$Z_{ij} = \mathbf{1}[\alpha_j \preceq \alpha_i], \quad C_i^u = \sum_j Z_{ij} \Pi_j^u.$$

The frozen Z has 129 ones. Its exact inverse $M = Z^{-1}$ has 65 nonzero entries in $\{-1, 0, 1\}$. This passage from cumulative values to atom values uses classical finite-poset Möbius inversion, not a technique introduced in this repository; see Rota (1964). The frozen acceptance contract checks both $MZ = I$ and $ZM = I$ as defense-in-depth; over the exact square rational matrices, either one-sided inverse identity already implies the other. Neither direction establishes event semantics: a consistently transposed or wrongly ordered pair can remain algebraically invertible.

Atom products and values are

$$R_i^u = \prod_{j=0}^{17} (Q_{\alpha_j}^u)^{M_{ij}}, \quad \Pi_i^u = \frac{1}{N} \ln R_i^u.$$

Because every R_i^u is a positive rational and $N > 0$, exact integer cross-multiplication decides

$$\Pi_i^u = 0 \iff R_i^u = 1, \quad \Pi_i^u > 0 \iff R_i^u > 1, \quad \Pi_i^u < 0 \iff R_i^u < 1.$$

A nonempty product may equal one. Exact zero is therefore not syntactic emptiness and not a floating tolerance. Magnitudes require a separate outward interval for $(1/N) \ln R$; interval overlap is not enough, and an interval touching zero does not decide exact sign.

Negative signed-net atoms are not merely hypothetical. The current bounded evidence matrix at `claims/SX-CERTIFIED-AVERAGED-PID3-001/evidence-matrix-v2.md` records $N = 3$, key 02+04, and exact net product $R^{\text{sx}} = 9/16$, hence $(1/3) \ln(9/16) < 0$. This is a bounded categorical witness, not a population or other-PID theorem. Clamping, taking an absolute value, or retaining only positive entries would change the declared object.

The audit registry is indexed by **18 three-source lattice positions** and contains **108 keyed scalar audit expressions**: 54 cumulative component expressions (18×3 for informative, misinformative, and signed net) plus 54 atom component expressions (18×3 for the same components). They are not 108 PID atoms, 108 lattice nodes, or 108 independent degrees of freedom. The number 166 belongs to the separate four-source SxPID4 carrier and is outside this three-source audit. The signed-net block is derived from the informative and misinformative blocks, while each atom block is a fixed linear transform of its cumulative block. The bounded binary qualification corpus has

$$\sum_{n=1}^5 \binom{n+15}{15} = \binom{21}{16} - 1 = 20,348$$

labeled count vectors and $20,348 \times 108 = 2,197,584$ coordinate verdicts. These are not all arbitrary-alphabet inputs and not 20,348 distinct probability laws; only 20,164 are primitive rational laws after division by the count-vector gcd.

12 Proposed proof-carrying semantic IR

The strongest prospective transfer is a two-view intermediate representation. It generalizes the existing certified-sxpid lowering pipeline to keyed SxPID3; it is not a new PID measure, and it is not yet a demonstrated engineering improvement.

12.1 View 1: semantic provenance graph

Each node retains a role-labelled key rather than only a number:

```
Total(N)
TargetMass(target=t)
EventMass(alpha, supported_row)
TargetEventMass(alpha, supported_row)
LocalRatio(component, alpha, supported_row)
Cumulative(component, alpha)
Atom(component, antichain_key)
```

The row count c_z is an exponent and provenance weight, not another multiplicative base. The graph records why a mass exists, which event produced it, which source/target values key it, and which coordinate consumes it. Numeric equality does not merge different semantic roles.

Required graph obligations are bidirectional:

1. **Forward completeness:** every mass and factor required by the frozen formulas appears exactly once under its semantic key.
2. **Backward soundness:** every stored node corresponds to one valid required object; there are no unreferenced or invented factors.
3. **Evaluation commutation:** evaluating the graph from the decoded count table yields the direct event/product specification for every key.
4. **Permutation equivariance:** renaming sources transports keys and values through one of the six source permutations without positional assumptions.

12.2 View 2: normalized exact arithmetic

Semantic keys are not a canonical arithmetic normal form. For example,

$$\frac{2 \cdot 6}{3 \cdot 4} = 1$$

despite a nonzero raw exponent vector over bases 2, 3, 4, 6. Therefore the acceptance route evaluates the semantic graph into a normalized positive rational, or compares its exact numerator/denominator cross-products after resource preflight. Structural cancellation may be a fast sufficient proof of zero; it is never the exact-zero authority.

Conversely, merging nodes only because their values happen to be equal would erase semantic mutation visibility. A target mass of six and an event mass of six are equal numbers but different claims. The two views solve different problems:

View	Preserves	Must not be used as
Semantic graph	roles, keys, provenance, completeness, mutation observability	canonical rational equality or exact sign authority
Normalized rational/cross-product	arithmetic cancellation, equality with one, strict sign	proof that factors came from the correct events or data

12.3 Baseline and adoption gate

audit/tools/certified-sxpid/src/exact.rs, extract.rs, and product.rs already provide the broad two-source pattern: canonical exact log arguments, count/event lowering, Möbius application, resource preflight, exact rational powers, and comparison with one. The first SxPID3 implementation must extend or independently reconstruct that simple direct route. The DAG is adopted only if it:


- closes a named Program C or D obligation without weakening semantics;
- is proved equivalent to direct reconstruction;
- reduces measured time, peak memory, certificate size, or review complexity on the exact corpus;

- keeps resource admission and public certificate semantics stable, or triggers a new claim revision when it does not; and
- survives the noncanonical-factor and semantic-role mutation suite.

Packed arrays and custom bit-level kernels remain deferred until profiling shows that this clearer route is inadequate.

13 Layered assurance architecture

The architecture separates claims that are often collapsed into the phrase “formally verified.”



No downstream green box repairs an upstream semantic error. The five closure programs are:

Program	Question	Main tools	Explicit nonclaim
A: semantic/combinatorial	Did the paper-defined event and lattice get transcribed correctly?	pinned source map, implementation-blind primary-source derivation, exact minimal witnesses	examples and citations are not a universal bridge
B: dual formal	Are carrier, order, matrix, event, product, and sign theorems derivable?	Lean concrete route; independently encoded SMT finite route	two kernels can share the same wrong source transcription
C: certified numerics	Do exact products and independent outward log intervals agree?	exact integers/rationals, MPFR/Arb producer, rational-series verifier	deterministic intervals are not confidence intervals
D: compiled refinement	Do the shipped Rust paths implement the keyed specification?	stable-key traces, serial/parallel/features/platform matrix	float agreement is not exact-product proof
E: adversarial replay	Are bytes, dependencies, mutations, first results, and custody complete?	content manifests, fresh replay, hostile fixtures, independent adjudication	hashes and reruns are not scientific truth

The complete target remains proposed/open for SX-CERTIFIED-AVERAGED-PID3-001. As recorded in decision record 2, Programs A and B have scoped partial results; Program C has a bounded exact sign/zero result; Program D has only a lexical routing observation; and Program E has source-bound local receipt and partial mutation evidence. None closes an end-to-end implication, and this report does not close the claim. Canonical-input obligation D1 remains an upstream requirement even though the frozen Program A–E “Current obligations” rows do not name it. Promotion requires every frozen obligation, including D1, as well as every Program closure condition; a five-row Program summary is not a complete obligation registry.

14 Formal methods as bounded instruments

Formal verification is most valuable when it narrows a specific uncertainty and least valuable when it becomes a ceremonial shield. Every evidence record should name its trusted computing base, independence dimensions, finite bounds, residual risks, prohibited inference, and the weakest plausible false claim it could still accept.

14.1 Lean

Lean can establish that a term inhabits the encoded proposition under its imports and axiom dependencies. It cannot by itself establish that the proposition is the intended paper theorem, that runtime bytes denote the intended data, or that an estimator is scientifically realistic.

Four notions must remain separate:

1. `hBV` : `BombieriVinogradov` is a theorem premise;
2. `propext`, `Classical.choice`, and `Quot.sound` are conventional proof-environment dependencies reported by an axiom audit;
3. intentional sorry placeholders in challenge modules are holes in a human comparator target, not holes in the default production theorem libraries; and
4. `noncomputable` marks an execution/extraction property, not an unproved or unsound theorem.

14.2 SMT

For finite masks, antichains, orders, and integer matrices, SMT is a strong second encoding. A theorem route proves unsatisfiability of the negation; each hostile mutation should become satisfiable with a retained model. unknown is no evidence. A wrong but self-consistent encoding can still yield `unsat`, so the SMT route starts from raw Boolean/set constraints rather than a generated Lean registry.

14.3 Exact arithmetic and directed intervals

Exact rational products eliminate rounding only for the encoded finite expression. They still rely on correct decoding, events, ordering, resource accounting, and big-integer execution. Directed log intervals rely on correct range reduction, tail bounds, endpoint rounding, and runtime behavior; they may legitimately abstain. They are deterministic enclosures, not statistical confidence intervals.

14.4 Exhaustive, mutation, metamorphic, and differential testing

- Exhaustive enumeration proves behavior only over its declared finite carrier.
- Mutation testing demonstrates sensitivity only to named faults.
- Metamorphic relations can be shared by multiple wrong implementations.
- Differential routes are useful only after common code, libraries, generated data, and semantic premises are mapped.
- A green CI run is one environment/revision observation. A warm-cache build is not a cold reproducibility receipt.

14.5 Hashes, councils, and paper graphs

A digest binds bytes up to collision risk; it does not prove truth, authorship, authenticity, meaning, or durable availability. A council generates attacks and alternative designs; shared models, prompts, sources, tools, or custody make it correlated. A paper graph or status badge is curation metadata until a same-revision checker binds entities, declarations, spans, imports, and evidence.

15 Attack matrix and required negative controls

The acceptance predicates must fail for the intended reason under at least the following classes.

Cut	Mutation or counterexample	Required observation
Decoder	duplicate member, noncanonical integer, negative-to-natural coercion, Unicode alias, reordered canonical bytes	rejection before event work
Event	OR/AND swap, OR within a mask, missing or moved target intersection	exact retained count witness disagrees
Weighting	uniform supported-row weight instead of exponent c_z	count-weighting witness rejects
Lattice	missing antichain, padded zero mask, reversed order, transposed Z , altered Möbius row	model or exact inverse/semantic reconstruction fails
Keying	positional (3)/(4) swap, source permutation without key transport	stable-key comparison rejects
Component	informative/misinformative swap, net multiplication instead of division, negative net clamp	exact quotient and reconstruction reject
Units	omit $1/N$, label bits as nats	identity and interval relation reject

Cut	Mutation or counterexample	Required observation
Factor IR	$(2 \cdot 6)/(3 \cdot 4) = 1$ with nonzero exponent vector	normalized rational says zero; structural-only route fails
Factor roles	equal-valued target/event masses merged, or c_z treated as a base	semantic graph completeness/soundness fails
Product	nonempty product one, factor/exponent mutation, huge projected power	exact comparison or preflight rejects
Interval	overlap substituted for containment, endpoint sign, mutated $\ln 2$, same-backend “independent” route	verifier rejects or abstains
Registry	partial 54-expression report, reordered identity, self-consistently resealed wrong registry	complete 108-key expression binding rejects
Corpus	replicated tables called distinct laws; missing or duplicate shard range when sharding is enabled	ordinal/coverage checker rejects
Refinement	specialized/general paths compared positionally or a feature path omitted	keyed execution matrix rejects
Interpretation	exact empirical sign called significance, calibration, or population truth	wording/schema gate rejects

16 Autoresearch without evidence laundering

Autoresearch is a candidate generator. It can be ambitious, adaptive, and computationally broad, but promotion rules depend on the question being asked.

16.1 Four distinct research loops

1. **Theorem discovery.** Models or search produce lemmas, proof outlines, witnesses, or counterexamples. Promotion requires a kernel-checked theorem plus source-correspondence review.
2. **Exact finite PID falsification.** Search produces canonical integer count tables that maximize disagreement, cancellation, resource use, or mutation sensitivity. Promotion requires independent exact reconstruction of the table and predicate.
3. **Estimator design.** Search proposes algorithms or hyperparameters. Promotion requires preregistered data-generating families, held-out and repeated confirmation, uncertainty, multiplicity control, robustness, and selection-aware reporting. An exact finite witness cannot substitute for these.
4. **Formal proof generation.** A model proposes terms or tactics. Kernel acceptance checks the formal proposition; human and source-mapping review still own intended meaning.

16.2 Promotion ledger

Before a run, record:

- objective and prohibited proxies;
- candidate grammar and resource budget;
- seed, environment, code revision, and allowed tools;
- stopping, withdrawal, and retention rules;
- judge identity and the candidate's access to it; and
- confirmation data or theorem surface held back from adaptive selection.

During search, retain every exposure that can influence later choices, not only winners. After search, freeze the candidate bytes before judging. The judge reconstructs from source input, emits one terminal result, and cannot be rewritten or queried adaptively by the candidate lane.

The research record keeps failures, near misses, exploit attempts, judge disagreements, and the first accepted result. A later clean rerun does not erase the selection history. Search frequency, score, model consensus, and repeated rediscovery are discovery signals, not evidence of the claim.

16.3 Council protocol

For current consequential designs, twenty mandatory core lenses and at least thirty additional artifact-specific lenses are applied in three passes. This publication uses fifty additional lenses, for seventy current lenses in total. The earlier twenty-lens transfer council remains a dated historical review; it is not the current coverage threshold:

1. independent pre-merge memos identify the strongest falsifier and most likely hidden shared cut;
2. an integrator reconstructs decisive facts from primary sources and synthesizes competing designs; and
3. an adversarial pass attacks the selected design, its evidence boundaries, and its wording.

The record preserves dissent, the strongest surviving objection, evidence that would reverse the decision, and independence limitations. There is no majority vote and no “assurance percentage.”

17 Repository durability and promotion

Simple migrations should follow one reusable state machine rather than produce a bespoke PDF. The human overview lives here; exact per-path facts belong in a strict JSON/TSV cleanup ledger. The 1 September ledgers linked below are dated pre-publication/pre-cleanup snapshots. The later post-publication custody receipt and its machine record record a separate, bounded observation of the subsequent publication and named cleanup actions. These records are chronological, not competing authorities: the first two freeze preconditions and grant no cleanup authority; the later pair records named outcomes and grants no future authority. None turns a dated snapshot into a live inventory. The historical bounded authorities are the primary-common-Git-directory retirement ledger and the sibling-registry retirement ledger. The primary ledger excludes sibling repositories. The sibling ledger covers ten other Git registries and twelve observed worktrees, but excludes the primary common Git directory and non-Git copies. Neither ledger authorizes deletion.

17.1 Why isolate work in branches and worktrees?

A separate branch gives an uncertain result a named history. A linked or detached worktree gives that history a separate checkout, index, and HEAD; project convention and build configuration can also give it a separate build root and bounded write location. This lets the project explore competing proofs, preserve a frozen claim revision, run a clean reproduction, or prepare an urgent correction without rewriting a dirty primary checkout. Git explicitly supports linked worktrees for concurrent branch checkouts and detached experimental work (Git worktree documentation).

Isolation has a cost. Linked worktrees share the repository's common object database and most refs; HEAD, the index, and some pseudorefs or per-worktree refs are worktree-specific. A branch can move. A clean commit omits dirty, untracked, and ignored bytes. A bundle contains selected refs and reachable Git objects, but not the worktree, index, stash, hooks, or repository configuration (Git bundle documentation). Separate agents can still share a wrong definition or stale checker. Therefore, a branch/worktree boundary is a write-isolation mechanism. It is not scientific independence, accepted status, durable storage, or proof that all useful work is captured.

Legitimate use	New risk	Control
Speculative theorem, estimator, counterexample, or document work	A mostly stale lane can still contain one unique useful result.	Compare paths, blobs, semantics, and negative evidence. Port the useful fragment; do not merge the stale tree wholesale.
Frozen proof or certificate revision	Checker, theorem, artifact, and toolchain pins can drift independently.	Bind them in one packet and require revision/digest mutations to fail.
Parallel agents with disjoint write scopes	Shared prompts, models, source maps, or files can create correlated errors and writer collisions.	Record ownership and shared dependencies. Stop writers before freeze. Treat council output as advisory.
Clean reproduction checkout	Ambient caches, dependencies, environment, or Git indirection can still influence the run.	Pin and inspect the trust surface. Repeat from a fresh route and in hosted CI on the exact commit.
Archive of a rejected, superseded, or restricted route	A side ref can disappear, remain inaccessible, or be mistaken for accepted mainline work.	Label its disposition, create a verified recovery object, and keep it outside executable/public authority surfaces.

17.2 What failed, what detected it, and what repaired it

The repository encountered a representative verification-chain drift: revision-4 Lean and SMT artifacts coexisted with checkers that still named revision 3. Pin drift alone does not establish whether the theorem is false or true. It establishes that the claimed replay chain is unusable from that checkout. The Lean path check and the SMT manifest check detected the mismatch before a digest comparison could establish artifact identity. The first repair decision is semantic: restore one complete revision-3 packet, migrate one complete revision-4 packet, or preserve and retire an obsolete lane. This repository selected revision 4, so its checker pins, artifacts, theorem roster, and hostile mutations moved together and were replayed in ordinary and optimized modes. Weakening the revision check, accepting whichever artifact exists, or resealing an expected digest without review would only hide the defect.

Other recurring failures and their repairs are:

Failure	Detection	Repair
Results split across clean commits, dirty files, untracked/ignored paths, and side refs	Common-Git-directory inventory plus per-path mode, size, digest, history, and process ownership	Classify every variant; integrate current accepted packets; preserve valuable failures and uncertain preimages separately
A stale bulk merge could restore old math or documentation over a corrected file	Semantic comparison against the current mainline object, followed by math/source/catalog/document gates	Port explicit paths or coherent claim packets only; regenerate derived state last
A local or side-branch success could be described as mainline completion	Direct remote-ref query and exact-SHA hosted-run inventory	Prove fast-forward ancestry, use an explicit expected-old lease, query the remote again, and require a new main-triggered run (Git push documentation)
Temporary build trees and caches consumed disk and obscured live ownership	Worktree registry, process table, cleanliness, and disk inventory	Remove only process-free, clean, superseded worktrees with Git-aware removal; delete reproducible caches after evidence preservation
A polished formal result could be read as probability or estimator assurance	Exact theorem, import, axiom, source-map, and nonclaim review	State the formal proposition and its assumptions; keep probability, estimator calibration, executable refinement, and application qualification as separate gates
A path census was reported without binding each count to its comparison target	Check the row's target identifier, supporting commit, basis, and exact/evolved/absent sum together	Keep distinct rows for the integration candidate and live remote main; never move one target's counts into another target's prose
A partial clone reported zero missing objects reachable from one observed HEAD	Inspect promisor configuration, partial-clone filter, pack markers, traversal root, and lazy-fetch setting	State only the observed-HEAD closure; replay a separately identified bundle in an isolated repository for its advertised ref; make no global-completeness claim
A commit unreachable from observed live refs was treated as if it lacked custody	Match the exact object to the advertised bundle inventory and isolated replay result	Record live-ref reachability and custody as separate fields; retain the registry only for remaining unadjudicated content, not for a custody gap that the evidence closes
An archived object identifier lost its leading hexadecimal digit	Enforce the full 40-hex SHA-1 grammar and cross-check the object, tree, and parent tuple	Restore the exact identifier and reject abbreviated or malformed identities before any ancestry or preservation conclusion

This failure ledger is part of the method. It explains why a green theorem can coexist with a red certification chain, why a branch can be mostly obsolete yet still be non-discardable, and why cleanup happens after integration and verification rather than before them. The dated worktree and branch preservation receipt provides the bounded repository case evidence: it records a lost temporary integration state, reconstruction limits, divergent-tip comparison, dirty-byte capture, bundle verification, and retirement preconditions. Its primary-repository facts at that dated review boundary are also bound by the machine-readable retirement ledger. The later multi-registry snapshot is bound by the machine-readable sibling-registry ledger. These records preserve target-specific comparisons, partial-clone limits, and exact custody boundaries. The later post-publication receipt adds the dated mainline/remote observation and explicitly names the state it does and does not cover. None promotes archived content, freshly replays an external registry, or replaces this reusable rule.

17.3 Formal verification and durable closure in one lane

For each formal lane, freeze the proposition, source-to-symbol map, assumptions, theorem roster, imports, axioms, toolchain, dependency commits, checker bytes, and expected artifacts. Run the kernel or solver. Run the checker in ordinary and optimized modes where applicable. Mutate the statement, revision pins, imports, expected artifacts, and acceptance predicates only through a preregistered suite of well-formed faults. Each fault names the obligation it violates and the causal rejection that must occur. Equivalent or stronger statements and irrelevant import changes are not required to fail. Add a separately developed route when it reduces a load-bearing shared cut. Then keep four downstream questions open unless they have their own evidence: does the formal statement match the intended PID definition; does the Rust/Python executable refine it; does the estimator behave under the declared population and sampling law; and is a consumer decision qualified?

The durability sequence is exact: freeze writers; inventory every visible and relevant ignored byte; classify each variant; commit or externally preserve unique material; bind bundle size and SHA-256; compare advertised refs with the frozen intended-ref inventory; run bundle verification and a clean recovery drill; integrate explicit paths from fresh main; regenerate Markdown/TeX/PDF and source-state derivatives; run applicable local gates; push the exact descendant with an expected-old lease; query remote main; require hosted success on that exact object; and

only then retire branches, worktrees, and caches. A post-removal inventory must show that no required object became unreachable.

FIGURE 3 - DURABLE PROMOTION

Retirement is the final closure gate

Seven ordered stages protect unique research bytes before a branch or worktree is retired.



The order prevents a false cleanup story:
discover every variant; freeze writers; classify its scientific and custody role;
preserve unique bytes; integrate exact paths; verify the relevant gates and remote state;
then retire only when recovery and a post-removal inventory pass.

NO DESTRUCTIVE RESET - NO PLAIN FORCE PUSH - NO GC BEFORE THIS GATE


Numbered stages and patterned fills supplement the palette; no scientific or accessibility assurance follows from this diagram.

Figure 3: Durable promotion state machine, part 1: the seven ordered stages from discovery through retirement.

FIGURE 4 - DURABLE PROMOTION

Availability, acceptance, and recovery are different claims

Archive-only material may be durable without becoming a mainline or scientific acceptance claim.



The storage and retirement diagrams are project-defined process guidance, not a release, custody, authenticity, or accessibility certificate.

Figure 4: Durable promotion state machine, part 2: the storage and promotion paths and retirement gate.

Discover -> Freeze -> Classify -> Preserve -> Integrate -> Verify -> Retire

17.4 Storage and trust ladder

Layer	Useful for	Not sufficient for
Scratch file / model context	rapid exploration	recovery, review, or durable identity
Dirty worktree	active local development	evidence that all bytes are tracked or backed up
Local commit/ref	recoverable local history	off-machine durability
Remote side/archive ref	recovery through the declared remote; off-host credit only when separate hosting is verified	accepted mainline status or scholarly permanence
Remote main	integrated project history	release identity or long-term external archive
Maintainer-approved project release with checked metadata; Git objects stay unsigned under repository policy	named versioned project milestone under the repository and service retention policies	immutability, independent scholarly deposit, or authenticity by itself
Scholarly or restricted archive; record its versioned locator or DOI when the archive supplies one	long-term citation/recovery under the identified repository and retention policy	automatic scientific correctness; a DOI alone is not byte custody
CI artifact	short-lived diagnostic exchange	archival storage; expiration is expected

The repository policy requires unsigned commits and tags and no AI attribution. Maintainer approval is a human governance decision, not a cryptographic Git signature.

17.5 Discover and freeze

Inventory every common Git directory, worktree, branch, tag, stash, checkpoint, unreachable object, dirty tracked path, untracked path, and relevant ignored path. Record HEAD, tree, merge base, mode, size, digest, and process

owner. Stop writers before taking a final manifest. Do not assume `git status` covers ignored research archives, caches that contain unique results, or alternate object stores.

17.6 Classify and preserve

Classify each byte variant as accepted science, open candidate, retained negative evidence, superseded-but-needed recovery state, restricted/local-only material, reproducible cache, or discardable temporary output. Secret-screen before copying. Environment files, credentials, device IDs, tokens, and private paths never enter a public snapshot.

Unique or uncertain material receives a named local ref bound to a recorded object ID and a verified, content-addressed second recovery copy or reviewed archive path before integration. Call that copy off-host only when its separate host/custodian and a retrieval drill establish the fact. A branch is not deleted because its name looks old; a worktree is not removed because its files resemble main.

17.7 Integrate

Start from fresh remote main, not from a dirty primary worktree. Port explicit paths or coherent claim packets. Reconcile changes by semantic role, not wholesale directory copy. Regenerate self-referential source states, receipts, and publication derivatives last. Keep commits small, unsigned, human-authored, and free of agent trailers.

For unpublished PID material, use separate reviewed successors for distinct scientific claims. Negative and withdrawn results remain reachable under explicit archive/evidence paths rather than being silently overwritten by accepted variants.

17.8 Verify and retire

Before push, require the relevant mathematical, formal, numerical, source-custody, build, test, lint, documentation, PDF, and mutation gates. After push, bind the exact remote head and hosted run identities. Local and remote main must resolve to the same intended commit without a non-fast-forward update or plain `--force`.

A branch or worktree is retirement-eligible only when:

1. no process owns it;
2. every dirty, untracked, and relevant ignored byte has an explicit disposition;
3. its commits are remote-main ancestors, or every unique payload has a durable successor and recovery ledger;
4. restricted material has a safe non-public home;
5. the remote-ref, ancestry, hosted-run, and recovery-drill checks pass; and
6. a post-removal inventory shows no newly unreachable required object.

Cache deletion and Git garbage collection come last, after a retention window and repeated object inventory. Remote archive branches remain until exact preimages have a main-reachable or verified external successor.

17.9 Review-boundary safety disposition (dated 2026-09-01)

At the review boundary, the dirty primary worktree contained substantial unpublished PID claim packets, scripts, reviews, and ignored material not proved redundant with remote main. It was therefore treated as a protected source, not a cleanup target. The SxPID3 packet used here was ported into this isolated lane with an exact preimage manifest, then received bounded presentation-only GitHub-Markdown normalization: legacy inline/display delimiters became $/$/$$, four `\operatorname` occurrences across three formulas became `\mathrm`, and one blockquoted implication was reflowed as ordinary prose plus display math. The ledger binds both preimage and current bytes. These mechanical edits were not intended to change mathematical or prose content, but the hashes alone do not prove semantic equivalence, and the preservation step does not adjudicate the packet. This section makes no cleanup or repository-wide integration claim for its review boundary. That conclusion remains true for the dated boundary. The later post-publication custody receipt records separately observed publication and named retirement actions; it neither rewrites this historical section nor becomes a live, global branch or filesystem authority.$

18 Implementation roadmap for SxPID3

The order below optimizes for semantic risk reduction, not visual progress. It is a prospective roadmap from the current decision-v2 boundary: it preserves partial bounded evidence without calling it an accepted certifier or silently re-opening the frozen revision-1 claim.

18.1 Phase 0: preserve and adjudicate the packet

- Preserve the frozen revision-1 packet, current decision-v2 record, and bounded negative results as distinct evidence objects.

- Keep the complete target proposed/open while the evidence-adjudication index remains the current pointer; do not edit the historical revision index to make later evidence appear contemporaneous.
- Pin the exact Makkeh--Gutknecht--Wibral source and construct a paper-equation-to-local-formula map.
- Decide whether the semantic DAG is internal revision-1 machinery or a claim-v2 acceptance object before changing an acceptance rule, schema, registry, or resource policy.

Stop condition: any unresolved source/event/unit/order ambiguity.

18.2 Phase 1: simple direct specification

- Implement the strict canonical decoder and direct event-count route.
- Reconstruct all 18 antichains, order entries, Z , M , 54 cumulative component expressions, 54 atom component expressions, and the 108 keyed scalar audit expressions without generated acceptance tables. Do not label those expressions as atoms or independent degrees of freedom.
- Add minimal exact witnesses for every semantic mutation and the noncanonical-factor example.

Promotion condition: implementation-disjoint, generated-answer-blind review and exact keyed results on hand-checkable fixtures.

18.3 Phase 2: dual formal semantics

- In Lean, define the concrete carrier, event membership, counts, order, zeta/Möbius relation, products, and sign reduction.
- Independently encode finite masks, antichains, order, all 324 zeta entries, inverse products, permutations, and event mutations for SMT.
- Freeze theorem/import/axiom rosters and run hostile substitutions.

Nonclaim: this does not close paper correspondence, log intervals, or Rust execution.

18.4 Phase 3: exact products and intervals

- Build a producer and dependency-disjoint verifier from canonical count bytes.
- Compare normalized exact rationals with one after aggregate/per-coordinate preflight.
- Produce directed intervals and require the verifier interval to be a subset of the producer interval.
- Make cancellation, timeout, precision exhaustion, storage failure, and resource refusal terminal, with no partial verified result, report, or shard root.

18.5 Phase 4: keyed Rust refinement

- Expose audit-only stable-key traces from specialized and general three-source routes.
- Compare serial/parallel, debug/release, features, platforms, package/binary identity, and Python exposure without positional assumptions.
- Keep binary64 tolerance results diagnostic; exact products and intervals remain the authority.

18.6 Phase 5: bounded corpus and optional shards

- Freeze a canonical ordinal for all 20,348 count vectors.
- Replay all 2,197,584 keyed-expression verdicts through the authoritative streaming route, plus permutations, relabelings, admitted replications, canonical gates, retained failures, cancellation, and resource refusals.
- If measured resource, replay, or recovery evidence justifies sharding, partition half-open ranges with no overlap or gap and prove the content-addressed root corresponds to the exact ordinal domain. Do not call it arbitrary-alphabet exhaustiveness.
- Keep shards as internal evidence packaging unless an adjudicated claim/route revision explicitly makes their schema or root acceptance-bearing.

18.7 Phase 6: promotion and external review

- Freeze source, tools, schemas, binaries, first-result records, negative evidence, and mutation receipts under one same-revision manifest.
- Obtain fresh acquisition and execution under separate custody.
- Adjudicate every frozen obligation, including D1, and every Program A--E closure condition. Any load-bearing open cut keeps the claim proposed or blocked.
- Publish one milestone MD/PDF and machine companions; do not create migration-specific narrative PDFs.

19 Evidence-card template

Every acceptance-bearing artifact should answer:

Field	Required content
Claim	Exact proposition and quantifiers supported by this item
Source semantics	What the originating object means in its own field
Bound	Input domain, count, precision, resource, platform, and revision limits
TCB	Kernels, runtimes, parsers, libraries, solvers, filesystem, and hardware assumptions
Independence vector	Different acquisition, custody, language, encoding, algorithm, backend, and reviewer dimensions
Shared cuts	Paper transcription, registry, parser, generator, arithmetic, data, or tooling dependencies
Defeaters	Conditions under which valid-looking evidence no longer supports the claim
Weak false claim accepted	The most plausible wrong proposition this route could still pass
Prohibited inference	What downstream statement must not cite this item
Reproduction	Exact inputs, command, outputs, hashes, and terminal status
Durability	Remote/release/archive locator, retention policy, and recovery test

20 Final recommendation

PrimeGapsLib offers a strong architectural lesson: a sophisticated search may remain outside the trusted argument if exact finite data are reconstructed by a smaller checker and connected to the final theorem through explicit bridges. It also offers a warning: elegant formal code, a large blueprint, a public build, and an interactive graph can coexist with mutable links, unexecuted comparators, warm-cache evidence, sparse process provenance, and no same-revision publication receipt.

For `pid-rs`, the best next step is not to imitate prime-gap mathematics or immediately build a packed kernel. It is to preserve the open `SxPID3` packet, close the paper-to-formula map, implement a simple generated-answer-blind exact route, and pilot a two-view proof-carrying semantic IR whose semantic and arithmetic obligations are separately testable. Only measured benefit justifies packing or internal sharding. Only closure of every frozen obligation, including `D1`, and every Program A--E closure condition can justify a certificate disposition. Nothing in that path upgrades KSG, continuous PID, population inference, or scientific interpretation without their own evidence.

The process recommendation is equally important: keep adaptive discovery broad, make promotion narrow, retain failures and selection history, expose shared cuts, generate one canonical human report from the same mathematical source, and move every valuable byte through an auditable storage-and-promotion ladder before retiring any branch or worktree.

21 Source-anchored claim register

The links below are pinned to the reviewed Git commit. They are locators for review, not a claim that a line range alone contains every transitive proof dependency.

The repository-wide search found six sorry placeholders, all in the two intentional challenge roots rather than the default theorem libraries. That fact does not turn challenge holes into production proof holes, nor does it establish that the challenge statements perfectly represent the paper. The comparator still needs an executed, pinned receipt before it earns verification credit.

Review claim	Primary anchor
Final public declaration is conditional	<code>PrimeGaps/Bounded246.lean</code> , lines 21--31
Stored certificate proves the cleared threshold corresponding to > 4	<code>PrimeGapsCert/Gap246/Witness.lean</code> , lines 35--46
4.004306 appears in a separate positivity-window theorem	<code>Sieve/PositivityWindow.lean</code> , lines 68--81
Explicit tuple card, diameter, and admissibility are checked	<code>H50.lean</code> , lines 31--50
Candidate generator is explicitly untrusted	<code>Emit/SourceGen.lean</code> , lines 10--15
Generator uses integer-to-natural conversion in labelled input decoding	<code>Emit/SourceGen.lean</code> , lines 69--89
Small kernel layer is isolated from the untrusted producer	<code>Kernel/Core.lean</code> , lines 9--14

Review claim	Primary anchor
Stored rows and the terminal inequality are reconstructed and packaged	Witness.lean, lines 19--46
Precomputed integer tables are converted to the explicit rational certificate	Certificate/Fast.lean, lines 4068--4134
The explicit rational certificate is bridged to the analytic L2 witness	ExplicitToAbstract.lean, lines 687--731
Human challenge and solution are separate roots	Challenge/Basic.lean, lines 8--48, Solution/Basic.lean, lines 1--3
Paper-tag URL and duplicate handling need hardening	PaperTag.lean, line 100 and lines 158--168
Reviewed public CI is a minimal tagged-action build	.github/workflows/lean.yml, lines 1--17
External PNT dependency is requested from mutable main	lakefile.toml, lines 20--23

22 Primary references

- Axiom Math, PrimeGapsLib at reviewed commit.
- Axiom Math, interactive bounded-prime-gaps paper.
- Axiom Math, formalization blueprint PDF.
- Makkeh, Gutknecht, and Wibral, *Introducing a differentiable measure of pointwise shared information*, Physical Review E 103, 032149 (2021), arXiv:2002.03356.
- Williams and Beer, *Nonnegative Decomposition of Multivariate Information* (2010), defining the distinct $I_{\min}$ redundancy functional.
- Bertschinger, Rauh, Olbrich, Jost, and Ay, *Quantifying Unique Information*, Entropy 16, 2161--2183 (2014), a basis for the distinct BROJA family.
- Ehrlich, Schick-Poland, Makkeh, Lanfermann, Wollstadt, and Wibral, *Partial Information Decomposition for Continuous Variables Based on Shared Exclusions: Analytical Formulation and Estimation*, Physical Review E 110, 014115 (2024).
- Gian-Carlo Rota, *On the Foundations of Combinatorial Theory I. Theory of Möbius Functions*, Zeitschrift für Wahrscheinlichkeitstheorie und Verwandte Gebiete 2, 340--368 (1964).
- pid-rs, MATHEMATICAL_PROBLEM_SOLVING_WORKFLOW.md, the existing repository-wide scientific and autoresearch process record.
- pid-rs, claims/SX-CERTIFIED-AVERAGED-PID3-001/, proposed revision-1 specification and open Programs A--E.